

Sundari Avanthikaa Srinivasan

AI / MACHINE LEARNING ENGINEER

Seattle, WA • avanthikaasrinivasan@gmail.com • linkedin.com/in/avanthikaa-srini • github.com/ventikaa • ventikaa.github.io

SUMMARY

AI/ML Engineer who takes models out of notebooks and into production: LLM, RAG, and agentic systems, plus the evaluation and MLOps discipline that keeps them reliable, accurate, and cost-aware at scale. M.S. Data Science candidate at the University of Washington and author of four peer-reviewed ML papers, with five internships shipping generative AI, computer vision, and full-stack ML to production.

CORE SKILLS

Languages: Python, SQL, TypeScript, JavaScript, Java, R

Machine Learning: PyTorch, TensorFlow, scikit-learn, XGBoost, deep learning, NLP, computer vision, SHAP

LLM & GenAI: LangChain, LangGraph, RAG, agentic workflows, prompt engineering, LLM evaluation, hallucination detection, Hugging Face

MLOps & Cloud: AWS (Lambda, S3), GCP, Docker, MLflow, FastAPI, CI/CD (GitHub Actions), Git, Linux

Data & Retrieval: pgvector, FAISS, embeddings, pandas, PySpark, Kafka, Airflow, Snowflake, MongoDB

EXPERIENCE

AI/ML Engineer Intern, **Transient.AI**

New York (Remote) · Jun 2026 – Present

- Build production LLM systems for capital markets: optimize RAG pipelines and agentic workflows (LangChain, LangGraph) over large financial-document corpora.
- Ship full-stack from retrieval to UI across a Python/FastAPI backend and a React/TypeScript frontend, deployed on AWS.

Software Engineering Intern, **EY (Ernst & Young)**

Chennai, India · Jan – Apr 2025

- Built a Spring Boot + MySQL platform serving 400+ employees with scalable REST APIs and automated reporting, cutting manual tracking overhead ~40%.

Digital & AI Intern, **TVS Motor Company**

Hosur, India · Sep – Nov 2024

- Deployed a real-time YOLO computer-vision pipeline at 83% accuracy for automated part detection, and engineered Python ETL that reduced quality-inspection time ~80%.

Generative AI Intern, **GreenPepper AI**

Remote · Jan – Apr 2024

- Built a production RAG system end-to-end (ingestion, pgvector retrieval, LLM synthesis, citations) serving 200+ users and cutting research-prep time ~40%.
- Designed LLM evaluation infrastructure: hallucination scoring, failure-mode taxonomy, adversarial red-teaming, and human-in-the-loop validation across four model families.

SELECTED PROJECTS

- **CLIP + FAISS Visual Search:** multimodal semantic image retrieval using CLIP embeddings indexed in FAISS for sub-100ms lookup, served via FastAPI. (*PyTorch, CLIP, FAISS*)
- **Healthcare Claims Fraud Detection:** anomaly-detection pipeline over 26M records deployed on AWS Lambda, with SHAP explanations for human review. (*XGBoost, scikit-learn, SHAP*)
- **LLM Jailbreak & Guardrail Study:** adversarial prompting pipelines and a failure taxonomy comparing where and why guardrails break across GPT, Claude, Llama 2, and Mistral. (*LLM evaluation*)

EDUCATION

University of Washington, M.S. Data Science

Seattle, WA · Sep 2025 – Jul 2027

SRM Institute of Science & Technology, B.Tech Computer Science (AI/ML)

2021 – 2025

PUBLICATIONS

Four peer-reviewed papers in machine learning across health informatics, computational neuroscience, and reinforcement learning (IEEE Xplore, Research Square, ICCETSP, IJSREM). Full list with links at ventikaa.github.io.